

CSE 6740: Homework 1

Mingzheng Michael Huang

September 14, 2023

Contents

1	Linear Regression [30 pts]	2
1.1	Derivation of $E[\hat{w}]$ [6 pts]	2
1.2	Derivation of $\text{Var}[\hat{w}]$ [6 pts]	2
1.3	Gaussian Distribution of \hat{w} [8 pts]	3
1.4	Weighted Linear Regression [10 pts]	4
2	Ridge Regression [15 pts]	5
3	Lasso Estimator	6
3.1	Equivalent Quadratic Function	6
3.2	Stochastic Gradient Descent Algorithm	6
3.3	Orthonormal X and SGD [10 pts]	6
4	Logistic Regression	8
4.1	Log-Odds as a Linear Function of X [6 pts]	8
4.2	Convexity of Logistic Loss [9 pts]	8
5	Programming: Recommendation System [40 pts]	9
5.1	Derivation of Update Formula [10 pts]	9
5.2	Regularized Objective Function [5 pts]	10
5.3	Python Implementation and Evaluation [15 pts]	11
5.4	Report [10 pts]	12
5.4.1	Performance (RMSE) on Training and Test Sets with Varied Low Rank	12
5.4.2	Observations with Varied Low Rank	12
5.4.3	Hyperparameter Selection	12

1 Linear Regression [30 pts]

1.1 Derivation of $E[\hat{w}]$ [6 pts]

Step 1: Starting with the normal equation, we have

$$\hat{w} = (X^T X)^{-1} X^T Y$$

Step 2: Taking the expectation on both sides

$$E[\hat{w}] = E[(X^T X)^{-1} X^T Y]$$

Step 3: Since X is fixed, it can be moved out of the expectation

$$E[\hat{w}] = (X^T X)^{-1} X^T E[Y]$$

We also know that $Y = Xw + \epsilon$, so $E[Y] = E[Xw + \epsilon] = Xw$:

$$E[\hat{w}] = (X^T X)^{-1} X^T Xw$$

Simplifying, we get:

$$E[\hat{w}] = w$$

Thus, \hat{w} is an unbiased estimator for w .

1.2 Derivation of $\text{Var}[\hat{w}]$ [6 pts]

The initial expression for the variance of \hat{w} is:

$$\text{Var}[\hat{w}] = (X^T X)^{-1} X^T (\sigma^2 I) X (X^T X)^{-1}$$

Step 1: Use the Identity Matrix

The term $\sigma^2 I$ is essentially scaling the identity matrix by σ^2 . Multiplying any matrix by the identity matrix leaves it unchanged. Thus, when we multiply $X^T (\sigma^2 I) X$, it's equivalent to scaling $X^T X$ by σ^2 .

Step 2: Simplify $\sigma^2 I$

Using the above reasoning, $X^T(\sigma^2 I)X = \sigma^2 X^T X$.

Step 3: Substitute into the Original Equation

Substituting this simplified term back into the original equation gives:

$$\text{Var}[\hat{w}] = (X^T X)^{-1} \sigma^2 X^T X (X^T X)^{-1}$$

Step 4: Cancel Out $X^T X$ Terms

Notice that $(X^T X)^{-1} X^T X (X^T X)^{-1}$ contains $(X^T X)^{-1} X^T X$, which will cancel out, leaving:

$$\text{Var}[\hat{w}] = \sigma^2 (X^T X)^{-1}$$

1.3 Gaussian Distribution of \hat{w} [8 pts]**Step 1: Expression for \hat{w}**

We start with the estimate \hat{w} in linear regression, which is given by:

$$\hat{w} = (X^T X)^{-1} X^T Y$$

Additionally, we have the true relationship between Y and X :

$$Y = Xw + \epsilon$$

Here, ϵ is the noise term, Gaussian distributed with zero mean and constant variance σ^2 .

Step 2: Substitute Y into \hat{w}

Next, we substitute Y into the equation for \hat{w} :

$$\hat{w} = (X^T X)^{-1} X^T (Xw + \epsilon)$$

Simplifying, we get:

$$\hat{w} = w + (X^T X)^{-1} X^T \epsilon$$

Step 3: Distribution of \hat{w}

1. **Linearity of Gaussian Distribution:** A linear combination of Gaussian variables is also Gaussian.
2. **ϵ is Gaussian:** ϵ follows a Gaussian distribution.

Thus, \hat{w} , being a linear combination of w (a constant) and ϵ (a Gaussian variable), must also be Gaussian.

Step 4: Mean and Variance of \hat{w}

We have derived that the expected value $E[\hat{w}]$ is w . We have also derived that $\text{Var}[\hat{w}] = \sigma^2(X^T X)^{-1}$.

So, under the white noise assumption, which stipulates that ϵ is Gaussian-distributed with zero mean and constant variance σ^2 , \hat{w} also follows a Gaussian distribution. Its mean is w and its variance is $\sigma^2(X^T X)^{-1}$.

1.4 Weighted Linear Regression [10 pts]

In weighted linear regression, the objective function $J(w)$ to be minimized is given by:

$$J(w) = (Y - Xw)^T \Sigma^{-1} (Y - Xw)$$

Step 1: Taking the Derivative

To find the value of w that minimizes $J(w)$, we take the derivative of $J(w)$ with respect to w and set it to zero. Mathematically, this is expressed as:

$$\frac{\partial J(w)}{\partial w} = 0$$

First, let's expand the objective function:

$$J(w) = Y^T \Sigma^{-1} Y - Y^T \Sigma^{-1} Xw - w^T X^T \Sigma^{-1} Y + w^T X^T \Sigma^{-1} Xw$$

Now, taking the derivative of $J(w)$ with respect to w yields:

$$\frac{\partial J(w)}{\partial w} = -2X^T \Sigma^{-1} Y + 2X^T \Sigma^{-1} Xw$$

Step 2: Finding the Estimator \hat{w}

Setting the derivative to zero gives:

$$-2X^T\Sigma^{-1}Y + 2X^T\Sigma^{-1}X\hat{w} = 0$$

Simplifying, we obtain:

$$\begin{aligned} X^T\Sigma^{-1}Y &= X^T\Sigma^{-1}X\hat{w} \\ \Rightarrow \hat{w} &= (X^T\Sigma^{-1}X)^{-1}X^T\Sigma^{-1}Y \end{aligned}$$

This is the estimator \hat{w} in weighted linear regression, accounting for the different variances σ_i^2 of each data point.

2 Ridge Regression [15 pts]

Step 1: Ridge Regression Estimate

The ridge regression estimate $\hat{\mathbf{w}}_{\text{ridge}}$ is:

$$\hat{\mathbf{w}}_{\text{ridge}} = (X^\top X + \lambda I)^{-1} X^\top Y$$

Step 2: Bayesian Posterior Mean Estimate

In a Bayesian framework, the mean of the posterior $\hat{\mathbf{w}}_{\text{posterior}}$ is:

$$\hat{\mathbf{w}}_{\text{posterior}} = (\tau^{-2}X^\top X + \sigma^{-2}I)^{-1}\tau^{-2}X^\top Y$$

Step 3: Comparison and Relationship

By setting λ such that $\lambda = \frac{\sigma^2}{\tau^2}$, we can establish:

$$\hat{\mathbf{w}}_{\text{ridge}} = \hat{\mathbf{w}}_{\text{posterior}}$$

Thus, the ridge regression estimate is the posterior distribution's mean under the Gaussian prior.

Step 4: Explicit Relation between λ , σ^2 , and τ^2

The explicit relationship is:

$$\lambda = \frac{\sigma^2}{\tau^2}$$

3 Lasso Estimator

3.1 Equivalent Quadratic Function

The LASSO regression problem can be formulated as

$$\min_{w \in \mathbb{R}^d} \sum_{i=1}^m (w^\top x_i - y_i)^2 \quad \text{subject to } \|w\|_1 \leq \lambda.$$

A transformation involves representing the original w as $w^+ - w^-$, where w^+ and w^- are the positive and negative components of w , respectively.

The new, equivalent function $Q(w^+, w^-)$ is

$$Q(w^+, w^-) = \sum_{i=1}^m ((w^+ - w^-)^\top x_i - y_i)^2 + \lambda (\|w^+\|_1 + \|w^-\|_1).$$

3.2 Stochastic Gradient Descent Algorithm

To calculate the gradient for a single data point (x_i, y_i) , we have:

$$\begin{aligned} \frac{\partial Q}{\partial w_j^+} &= 2((w^+ - w^-)^\top x_i - y_i) \cdot x_{ij} + \lambda, \\ \frac{\partial Q}{\partial w_j^-} &= -2((w^+ - w^-)^\top x_i - y_i) \cdot x_{ij} + \lambda. \end{aligned}$$

After obtaining the gradient, the update rules for w_j^+ and w_j^- are:

$$\begin{aligned} w_{\text{new}}^+ &= w_{\text{old}}^+ - \alpha \frac{\partial Q}{\partial w_j^+}, \\ w_{\text{new}}^- &= w_{\text{old}}^- - \alpha \frac{\partial Q}{\partial w_j^-}. \end{aligned}$$

3.3 Orthonormal X and SGD [10 pts]

Explanation

First, while the orthonormality simplifies the optimization problem, it doesn't ensure that SGD will find the sparse solution w with k non-zero elements. Second, factors like learning rate and initial conditions play a role in the

convergence of SGD, but they don't necessarily guide it toward a sparse solution.

Furthermore, the inherent randomness in stochastic methods like SGD introduces an additional layer of uncertainty in finding the exact sparse solution. While LASSO's ℓ_1 -norm constraint does promote sparsity, the primary aim of SGD is to minimize the loss function. This could theoretically allow it to converge to w , but there's no practical guarantee.

In summary, despite the favorable conditions and the sparsity-promoting ℓ_1 -norm, SGD is not tailored to find sparse solutions. Therefore, there's no guarantee it will converge arbitrarily close to w .

4 Logistic Regression

4.1 Log-Odds as a Linear Function of X [6 pts]

The log-odds of success, or the logit, is defined as:

$$\ln \left(\frac{P(Y = 1|X = x)}{P(Y = 0|X = x)} \right)$$

Given that

$$P(Y = 1|X = x) = \frac{\exp(w_0 + w^T x)}{1 + \exp(w_0 + w^T x)}$$

and

$$P(Y = 0|X = x) = 1 - P(Y = 1|X = x) = \frac{1}{1 + \exp(w_0 + w^T x)}$$

Substituting these into the log-odds equation yields:

$$\ln \left(\frac{\exp(w_0 + w^T x)}{1 + \exp(w_0 + w^T x)} \times \frac{1 + \exp(w_0 + w^T x)}{1} \right)$$

Simplifying, we get:

$$\ln(\exp(w_0 + w^T x)) = w_0 + w^T x$$

Thus, the log-odds of success is a linear function of X .

4.2 Convexity of Logistic Loss [9 pts]

To show that the logistic loss function $L(z) = \log(1 + \exp(-z))$ is convex, we take its second derivative.

First derivative $L'(z)$:

$$L'(z) = \frac{-\exp(-z)}{1 + \exp(-z)}$$

Second derivative $L''(z)$:

$$L''(z) = \frac{\exp(-z)}{(1 + \exp(-z))^2}$$

Since $L''(z)$ is always positive, $L(z)$ is a convex function.

5 Programming: Recommendation System [40 pts]

5.1 Derivation of Update Formula [10 pts]

Partial Derivatives

To minimize $E(U, V)$, we employ the gradient descent optimization technique, which necessitates the computation of partial derivatives of $E(U, V)$ with respect to each element in U and V .

Partial Derivative with respect to $U_{v,k}$

The partial derivative of $E(U, V)$ concerning $U_{v,k}$ can be computed as follows:

$$\frac{\partial E(U, V)}{\partial U_{v,k}} = -2 \sum_{i=1}^m (M_{v,i} - U_v V_i^\top) V_{i,k}$$

Partial Derivative with respect to $V_{j,k}$

Similarly, the partial derivative of $E(U, V)$ concerning $V_{j,k}$ is:

$$\frac{\partial E(U, V)}{\partial V_{j,k}} = -2 \sum_{u=1}^n (M_{u,j} - U_u V_j^\top) U_{u,k}$$

Update Rule for U

$$U_{v,k} \leftarrow U_{v,k} + 2\mu \sum_{i=1}^m (M_{v,i} - U_v V_i^\top) V_{i,k}$$

Update Rule for V

$$V_{j,k} \leftarrow V_{j,k} + 2\mu \sum_{u=1}^n (M_{u,j} - U_u V_j^\top) U_{u,k}$$

where μ is the learning rate, a hyperparameter that governs the magnitude of each update step.

5.2 Regularized Objective Function [5 pts]

To improve the robustness of the matrix-factorization model against overfitting, we introduce regularization terms into the objective function $E(U, V)$. The new regularized objective function is as follows:

$$E(U, V) = \sum_{u=1}^n \sum_{i=1}^m \left(M_{u,i} - \sum_{k=1}^d U_{u,k} V_{i,k} \right)^2 + \lambda \sum_{u,k} U_{u,k}^2 + \lambda \sum_{i,k} V_{i,k}^2$$

Derivation of Partial Derivatives

To minimize this objective function, we need to compute the gradient with respect to each element in the matrices U and V . The partial derivatives with respect to $U_{v,k}$ and $V_{j,k}$ are derived as follows:

Partial Derivative with respect to $U_{v,k}$

$$\frac{\partial E(U, V)}{\partial U_{v,k}} = -2 \sum_{i=1}^m \left(M_{v,i} - \sum_{k=1}^d U_{v,k} V_{i,k} \right) V_{i,k} + 2\lambda U_{v,k}$$

Partial Derivative with respect to $V_{j,k}$

$$\frac{\partial E(U, V)}{\partial V_{j,k}} = -2 \sum_{u=1}^n \left(M_{u,j} - \sum_{k=1}^d U_{u,k} V_{j,k} \right) U_{u,k} + 2\lambda V_{j,k}$$

Update Rules via Gradient Descent

Utilizing these partial derivatives, the update rules for U and V using gradient descent with step size μ are as follows:

Update Rule for U

$$U_{v,k} \leftarrow U_{v,k} - \mu \left(-2 \sum_{i=1}^m \left(M_{v,i} - \sum_{k=1}^d U_{v,k} V_{i,k} \right) V_{i,k} + 2\lambda U_{v,k} \right)$$

Update Rule for V

$$V_{j,k} \leftarrow V_{j,k} - \mu \left(-2 \sum_{u=1}^n \left(M_{u,j} - \sum_{k=1}^d U_{u,k} V_{j,k} \right) U_{u,k} + 2\lambda V_{j,k} \right)$$

These update rules incorporate ℓ_2 regularization to mitigate the risk of overfitting by penalizing large values in U and V .

5.3 Python Implementation and Evaluation [15 pts]

```
1 import numpy as np
2
3 def my_recommender(rate_mat, lr, with_reg):
4
5     max_iter = 250
6     learning_rate = 0.0001
7     reg_coef = 2 if with_reg else 0
8
9     n_user, n_item = rate_mat.shape
10
11     U = np.random.rand(n_user, lr)
12     V = np.random.rand(n_item, lr)
13
14     mask = np.ma.masked_where(rate_mat != 0, rate_mat).mask
15
16     for i in range(max_iter):
17         U -= learning_rate * 2 * np.matmul(
18             np.multiply(np.matmul(U, np.transpose(V)) - rate_mat, mask),
19             V -= learning_rate * 2 * np.matmul(
20                 np.multiply(np.transpose(np.matmul(U, np.transpose(V)) - rate
21
22         U -= learning_rate * 2 * reg_coef * U
23         V -= learning_rate * 2 * reg_coef * V
24
25     return U, V
```

5.4 Report [10 pts]

5.4.1 Performance (RMSE) on Training and Test Sets with Varied Low Rank

Method	Low Rank	Training RMSE	Test RMSE	Time (s)
SVD-noReg	1	0.9319	0.9609	14.28
SVD-withReg	1	0.9368	0.9657	13.05
SVD-noReg	3	0.9175	0.9556	18.32
SVD-withReg	3	0.9231	0.9596	19.60
SVD-noReg	5	0.9072	0.9536	22.44
SVD-withReg	5	0.9127	0.9558	25.66
SVD-noReg	7	0.9003	0.9564	21.48
SVD-withReg	7	0.9046	0.9563	16.13
SVD-noReg	9	0.8942	0.9569	15.96
SVD-withReg	9	0.8968	0.9548	20.76

5.4.2 Observations with Varied Low Rank

An evident trend is observed where the RMSE decreases as the `lowRank` value increases, indicating improved model performance. However, the marginal gains in RMSE reduction appear to taper off as the rank grows, suggesting a point of diminishing returns beyond which increasing the rank doesn't contribute significantly to performance improvement.

5.4.3 Hyperparameter Selection

- Learning Rate (μ): Chosen as 0.0001 after experimenting with multiple values to ensure quick convergence without oscillations.
- Regularization Coefficient (λ): Set to 2 when regularization is applied. This value was empirically determined to safeguard against overfitting without compromising the predictive accuracy.